

# Professional Profile – Siddharth Chamoli

## Personal Overview

Siddharth Chamoli is an AI and Machine Learning Engineer with over 6 years of professional experience in data science, intelligent automation, and large-scale analytics systems. He has worked extensively with enterprise clients in banking, telecom, and technology domains.

## Current Role

Siddharth currently works with IRIS Software on banking and financial services projects. His work focuses on building production-grade AI systems, MLOps pipelines, and cloud-native ML services primarily using AWS infrastructure.

## Core Technical Skills

- Python, FastAPI, REST API design
- Machine Learning and Deep Learning
- Natural Language Processing (NLP)
- Retrieval-Augmented Generation (RAG) systems
- FAISS and vector databases
- OpenAI and Hugging Face models
- Docker and Kubernetes for ML services
- AWS (S3, EC2, ECR, EKS, CloudWatch)
- Git, GitHub, CI/CD pipelines

## Notable Projects

- Built an end-to-end RAG system with document ingestion, chunking, embeddings, FAISS indexing, and grounded LLM responses with source attribution.
- Developed intelligent document processing (IDP) pipelines for banking and insurance clients.
- Designed AI-driven analytics and competitor intelligence systems using large language models.

## Career Goals

Siddharth aims to become a Senior AI Engineer and AI SaaS builder capable of designing, deploying, and monetizing production-ready AI products across web and mobile platforms. He is particularly focused on combining GenAI with strong MLOps practices.

This document is intended for testing document ingestion, chunking, indexing, and querying within Retrieval-Augmented Generation (RAG) systems.